

Stochastics - Homework 5

Moira Lampe

NOTE: I was not able to finish in time. Please send me an email to moira.lampe@outlook.com, so I can send you the completed version. Thank you!

Point Processes

① Maximum likelihood for Markov Chain

Let's imagine a cognitive context where we are interested in sleeping patterns. Every 5 minutes, a check is performed on the participants & we note their state.

We want to model this by a (discrete) Markov chain with 3 states : "Sleeping", "(fully) awake", "drowsy".

- a) Plot the diagrams & parametrize the two following (informal) models with the least possible amount of parameters

• First case

① If $x=0$, we have three possibilities : stay at 0
go to 1



② If $x=1$, we have three possibilities : stay at 1
go to 0
go to 2

③ If $x=2$, we have three possibilities : stay at 2
go to 1

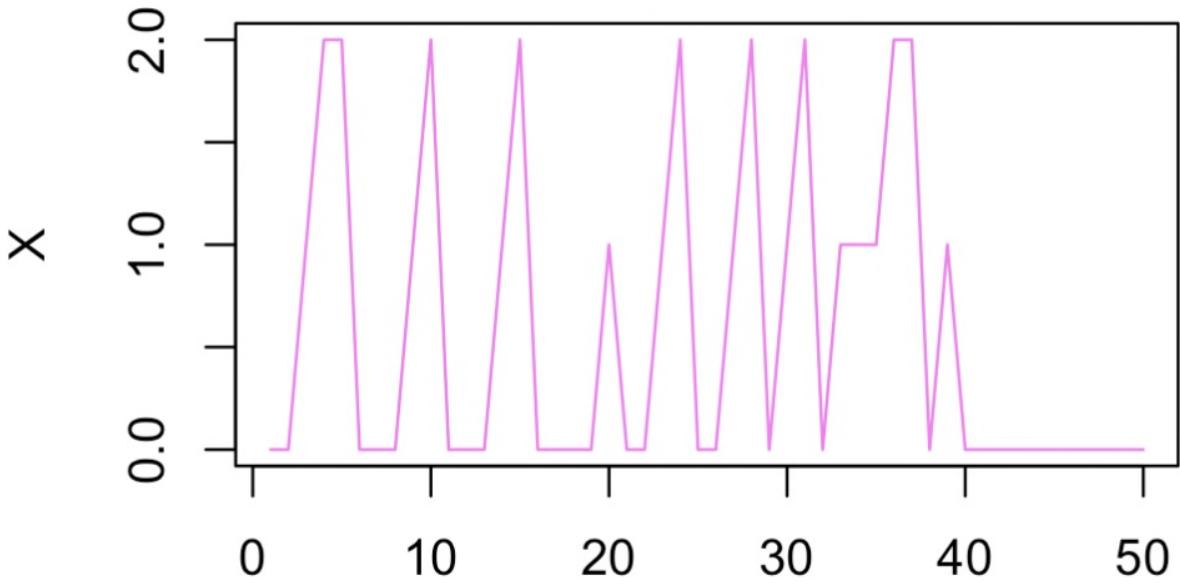
```

#First case

Y = rbinom(50,1, 0.3)
print(Y)
Z = rbinom(50,1, 0.7)
print(Z)
W = rbinom(50,1, 0.4)
print(W)
U = rbinom(50,1, 0.5)
print(U)
X <- c()
for (i in 2:50){
  X[1] = 0 #initializing state
  if (X[i-1] == 0){
    X[i] <- X[i-1]*(1 - Z[i])+(1 - X[i-1])*Y[i]
  }
  else if (X[i-1] == 1){ #we are in state drowsy
    #if we are in state drowsy, we can go back to sleep or awake, so we can go to 0 or to 1, and we can stay in :
    X[i] <-X[i-1]*(1 + W[i]) - X[i - 1]*Y[i]
  }
  else if (X[i-1] == 2){
    X[i] <- X[i-1]*(1 - U[i]) #if U is 0, we stay in 2, if U is 1, we go back to 1
  }
}
plot(1:50, X, type = 'l', main = 'First Markov Chain', col = 'violet')

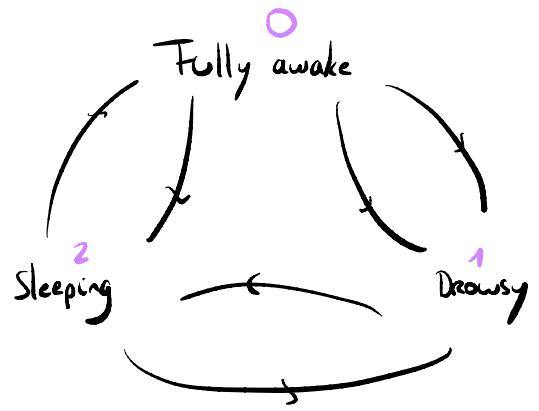
```

First Markov Chain



1:50

• Second case



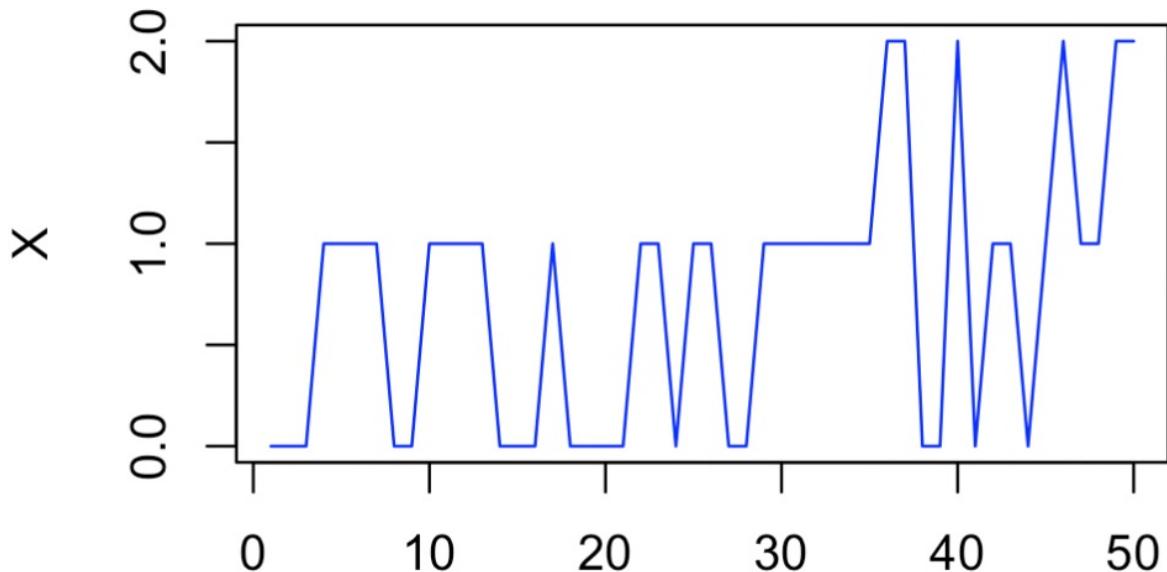
- ① If $X=0$, we have three possibilities : stay at 0
go to 1
go to 2
- ② If $X=1$, we have three possibilities : stay at 1
go to 0
go to 2
- ③ If $X=2$, we have three possibilities : stay at 2
go to 0
go to 1

#Second case

```

Y = rbinom(50,1, 0.3)
print(Y)
Z = rbinom(50,1, 0.6)
print(Z)
W = rbinom(50,1, 0.2)
print(W)
U = rbinom(50,1, 0.8)
print(U)
P = rbinom(50,1, 0.2)
print(P)
G = rbinom(50,1, 0.5)
print(G)
X <- c()
for (i in 2:50){
  X[i] = 0 #initializing state
  if (X[i-1] == 0){
    X[i] <- Y[i] + P[i]
  }
  else if (X[i-1] == 1){ #we are in state drowsy
    #if we are in state drowsy, we can go back to sleep or awake, so we can go to 0 or to 1, and we can stay in 1
    X[i] <- W[i] + Z[i]
  }
  else if (X[i-1] == 2){
    X[i] <- U[i] + G[i] #if U is 0, we stay in 2, if U is 1, we go back to 1
  }
}
plot(1:50, X, type = 'l', main = 'Second Markov Chain', col = 'blue')
plot(1:50, X, main = 'First Markov Chain', col = 'red')
  
```

Second Markov Chain



1:50

b) Make a function in R which simulates the first model if you give in the entries the state x_0 , the parameters of the model (only the ones parametrizing the transitions) and the length of the chain to simulate n .

```
#First case
MarkovChain_1 = function(p_01, p_10, p_12, p_21, n, X_0){
  Y = rbinom(n,1, p_01)
  print(Y)
  Z = rbinom(n,1, p_10)
  print(Z)
  W = rbinom(n,1, p_12)
  print(W)
  U = rbinom(n,1, p_21)
  print(U)
  X <- c()
  for (i in 2:n){
    X[1] = X_0 #initializing state
    if (X[i-1] == 0){
      X[i] <- Y[i]
    }
    else if (X[i-1] == 1){ #we are in state drowsy
      #if we are in state drowsy, we can go back to sleep or awake, so we can go to 0 or to 1, and we can stay in 1
      X[i] <- Z[i] + W[i]
    }
    else if (X[i-1] == 2){
      X[i] <- U[i] + 1 #if U is 0, we stay in 2, if U is 1, we go back to 1
    }
  }
  return(X)
}

MarkovChain_1(0.5, 0.5, 0.5, 0.5, 10, 0)
print(X)
```

#Second Case

```
MarkovChain_2 = function(p_01, p_10, p_12, p_21, p_02, p_20, n, X_0){  
  Y = rbinom(n,1, p_01)  
  print(Y)  
  Z = rbinom(n,1, p_10)  
  print(Z)  
  W = rbinom(n,1, p_12)  
  print(W)  
  U = rbinom(n,1, p_21)  
  print(U)  
  P = rbinom(n,1, p_02)  
  print(P)  
  G = rbinom(n,1, p_20)  
  print(G)  
  X <- c()  
  for (i in 2:n){  
    X[1] = X_0 #initializing state  
    if (X[i-1] == 0){  
      X[i] <- Y[i] + P[i]  
    }  
    else if (X[i-1] == 1){ #we are in state drowsy  
      #if we are in state drowsy, we can go back to sleep or awake, so we can go to 0 or to 1, and we can stay in 1  
      X[i] <- W[i] + Z[i]  
    }  
    else if (X[i-1] == 2){  
      X[i] <- U[i] + G[i] #if U is 0, we stay in 2, if U is 1, we go back to 1  
    }  
  }  
  return(X)  
}  
  
MarkovChain_2(0.2, 0.9, 0.2, 0.1, 0.1, 0.8, 30, 0)
```

c) Write the (log) likelihood of this problem and give the maximum likelihood estimator of p_{0-s} , the probability to fall asleep when you are drowsy. Would it be different if we were in the second model?

For a given model, $X \sim P_\Theta, g_\Theta(x) = P_\Theta(X=x)$, the likelihood will be computed as follows:

$$\Theta \rightarrow L_\Theta(x) = P_\Theta(X=x)|_{x=x}$$

The MLE is $\hat{\Theta} = \underset{\Theta \in \Theta}{\operatorname{argmax}} L_\Theta(x) = \underset{\Theta \in \Theta}{\operatorname{argmax}} \underbrace{L_\Theta(x)}_{l=\text{log likelihood}}$

\Rightarrow We need Θ to be the minimal possible set of parameters

\hookrightarrow as here we are only interested in p_{0-s} , so going from state 1 to 0: $p_{1 \rightarrow 0}$, we can say that we have full knowledge of the setup if we know $P = P(X_0=0)$ and $p_{0 \rightarrow 1}$ & $p_{1 \rightarrow 0}$

Our parameter space here is: $\Theta = (P, p_{0 \rightarrow 1}, p_{1 \rightarrow 0})$

For a given realisation x_0, \dots, x_n :

$$\begin{aligned} P_\Theta((X_0, \dots, X_n)) &= ((x_0, \dots, x_n)) = P(X_0=x_0, \dots, X_n=x_n) \\ &= P_\Theta(X_0=x_0, \dots, X_{n-1}=x_{n-1}) \cdot P_\Theta(X_n=x_n | \text{all the points from } 0 \text{ to } n-1) \\ &= P_\Theta(X_0=x_0) \cdot \dots \cdot P_\Theta(X_n=x_n | X_{n-1}=x_{n-1}) \\ &= p^{x_0}(1-p)^{1-x_0} \cdot p_0^{x_1} \cdot p_{1 \rightarrow 0}^{x_2} \cdot p_{0 \rightarrow 1}^{x_3} \cdot p_{1 \rightarrow 0}^{x_4} \end{aligned}$$

With $n_{i \rightarrow j}$ the number of transition from i to j seen in the sequence x_0, \dots, x_n .

Likelihoods:

So the likelihood sequence X_0, \dots, X_n with counts N_{i-j} is the number of times a transition $i \rightarrow j$ occurs:

$$L_\theta(X) = \rho^{X_0} (1-\rho)^{1-X_0} \cdot p_{0 \rightarrow 1}^{N_{0 \rightarrow 1}} \cdot (1-p_{0 \rightarrow 1})^{N_{1 \rightarrow 0}} \cdot p_{1 \rightarrow 0}^{N_{1 \rightarrow 0}} (1-p_{1 \rightarrow 0})^{N_{0 \rightarrow 1}}$$

So the log likelihood is:

$$\begin{aligned} \ell_\theta(x) &= X_0 \log(\rho) + (1-X_0) \log(1-\rho) + N_{1 \rightarrow 0} \log(p_{1 \rightarrow 0}) \\ &\quad + N_{1 \rightarrow 1} \log(1-p_{1 \rightarrow 0}) + N_{0 \rightarrow 0} \log(1-p_{0 \rightarrow 1}) \\ &\quad + N_{0 \rightarrow 1} \log(p_{0 \rightarrow 1}) \end{aligned}$$

To find the MLE of the transitions:

$$\frac{d \ell_\theta(x)}{d p_{1 \rightarrow 0}} = \frac{N_{1 \rightarrow 0}}{p_{1 \rightarrow 0}} - \frac{N_{1 \rightarrow 1}}{1-p_{1 \rightarrow 0}}$$

The derivative vanishes when $\frac{N_{1 \rightarrow 0}}{p_{1 \rightarrow 0}} = \frac{N_{1 \rightarrow 1}}{1-p_{1 \rightarrow 0}}$

$$\text{that is } N_{1 \rightarrow 0} = (N_{1 \rightarrow 1} + N_{1 \rightarrow 0}) p_{1 \rightarrow 0} \text{ that is } p_{1 \rightarrow 0} = \frac{N_{1 \rightarrow 0}}{N_{1 \rightarrow 1} + N_{1 \rightarrow 0}}$$

Hence:

$$\hat{p}_{1 \rightarrow 0} = \frac{N_{1 \rightarrow 0}}{N_{1 \rightarrow 1} + N_{1 \rightarrow 0}}$$

The number of times "1 to 0" occurred divided by the number of times 1 was seen.

$$\frac{N_{i \rightarrow j}}{\#\{X_k=i\}}$$

\Rightarrow This estimator is consistent if the chain is recurrent. Recurrence means that the chain will visit the transition an infinite number of times when $n \rightarrow \infty$

It would be different if one of the models were not recurrent.
State i is recurrent if "starting from i , and from wherever you can go, there is a way of returning to i ".

Transient: at least one state where you cargo & then there will be no path back to it.
→ as there are no transient states in both models, it will not be different.

d) Show on simulation that this estimator converges towards the true probability when n tends to infinity.

Let's say $p_{10} = 0.5$

By simulating with large number n , we can see that $\hat{p}_{10} \rightarrow p_{10}$

```
amount_drowsy = function(X, n){  
    ones <- c()  
    for (k in 1:n){  
        if (X[k] == 1){  
            ones = c(ones, +1)  
            print(ones)  
        }  
    }  
    return(n_ones = length(ones))  
}  
  
drowsy_to_sleep = function(X, n){  
    one_to_zero <- c()  
    for (j in 1:n){  
        if (X[j-1] == 0 && X[j] == 1){  
            one_to_zero = c(one_to_zero, +1)  
            print(one_to_zero)  
        }  
    }  
    return(n_10 = length(one_to_zero))  
}  
  
p_10_hat = function(n_10, n_ones){  
    return(n_10/n_ones)  
}
```

c) How would you decide if we are in the first or second model when you have the observations?

→ See figures

② Maximum likelihood for Poisson Processes

We are observing an inhomogeneous Poisson Process with intensity $\lambda(t)$ on $[0, T_{\max}]$.

We want to estimate $\lambda(t)$ by a piecewise constant function on a regular partition which means that the candidate intensities λ_i are, for a certain integer d giving the number of bins, of the form:

$$\lambda_i = \sum_{i=1}^d a_i^d \mathbb{1}_{b_{i-1}^d < t \leq b_i^d}, \quad \text{where } b_i^d = \frac{iT_{\max}}{d}$$

- a) When $d=1$, write the likelihood of the process Δ provide & MLE of a_1 . Implement it & show on simulation that when T_{\max} grows, the estimation is converging when the process is homogeneous.

Likelihood of N is:

$$L_\lambda(N) = \prod_{t \in N} \lambda(t) \cdot e^{-\int_0^{T_{\max}} \lambda(t) dt}$$

The log likelihood of N is:

$$l_\lambda(N) = \sum_{t \in N} \log(\lambda(t)) - \int_0^{T_{\max}} \lambda(t) dt$$

$$= \int_0^{T_{\max}} \log(\lambda(t)) dN_t - \int_0^{T_{\max}} \lambda(t) dt$$

$$l_{[a,b]}(N) = \log(n \cdot a) N_{[0, \frac{T_{\max}}{2}]} + \log(n \cdot b) N_{[\frac{T_{\max}}{2}, T_{\max}]} - na \frac{T_{\max}}{2} - nb \frac{T_{\max}}{2}$$

$$\hat{a} = \frac{N_{[0, \frac{T_{\max}}{2}]}}{n \frac{T_{\max}}{2}}$$

$\frac{d}{da}$

$$\hat{b} = \frac{N_{[\frac{T_{\max}}{2}, T_{\max}]}}{n \frac{T_{\max}}{2}}$$

$$\begin{aligned}
 \text{Because } d=1, \text{ we have that } \lambda(t) &= \sum_{i=1}^d a_i^1 \mathbb{1}_{b_{i-1}^1 < t \leq b_i^1} \\
 &= \sum_{i=1}^1 a_1^1 \mathbb{1}_{b_{1-1}^1 < t \leq b_1^1} \\
 &= a_1^1 \mathbb{1}_{0 < t < T_{\max}}
 \end{aligned}$$

If we fill in from above

$$= \int_0^{T_{\max}} \log(\lambda(t)) dN_t - \int_0^{T_{\max}} \lambda(t) dt$$

We have that

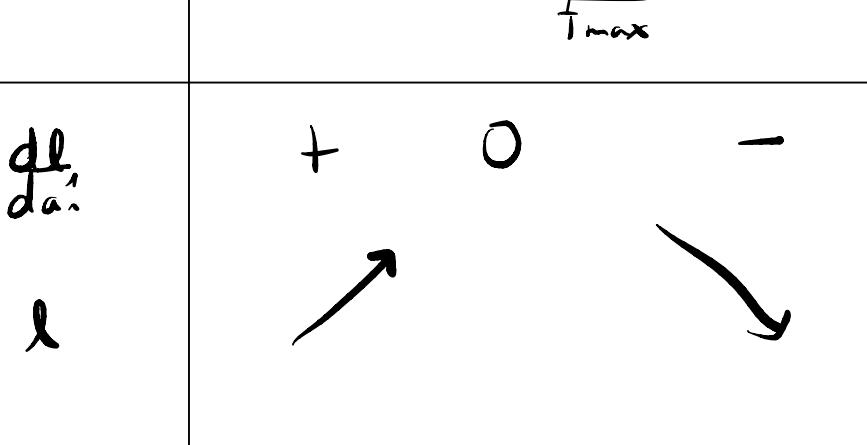
$$l = \log a_1^1 N_{[0, T_{\max}]} - a_1^1 T_{\max}$$

$$\text{Doing our derivation : } \frac{dl}{da_1^1} = \frac{1}{a_1^1} N_{[0, T_{\max}]} - T_{\max}$$

Hence, we find that our MLE is:

$$\hat{a}_1^1 = \frac{N_{[0, T_{\max}]}}{T_{\max}}$$

$$\frac{N_{[0, T_{\max}]}}{T_{\max}}$$



I was not able to prove it...

```
#function for homogeneous poisson process
Homo_Poisson = function(Tmax, lambda)
{
  t = 0 #t starts out with zero
  N = numeric()
  res = c()
  while(t < Tmax){ #we want the process to go on until the cumulative sum of t reaches T, which is 1 here
    u = runif(1) #as long as the above is true, we create u, which is a uniform random variable
    t = t - log(u)/lambda #then our new t will be this #it is the same thing as doing it separately and then taking the cumulative sum of it.
    N = c(N, t)
  }
  res = c(N, "The a_hat is:", sum(N/Tmax), "the sum of N is:", sum(N))
  return(res)
}

Homo_Poisson(10, lambda = 3)
#this is not working, it is definitely not converging when Tmax grows...
```

```
#Let's try a different way of simulating a poisson process:
#Let's try with exponentially distributed interarrival times.
#We simulate the arrival times until the maximum time horizon is achieved.
```

```
My_Poisson = function(Tmax, rate){
  Y = c()
  arrivals <- rexp(50, rate)
  while (arrivals[length(arrivals)] < Tmax){
    Y = c(Y, arrivals)
  }
  return(Y)
}
```

```
My_Poisson(1, 0.4)
print(Y)
```

```
#I thought I could simulate it differently but somehow this one also doesn't work
```

b) For a more general d write the corresponding likelihood & give the MLE of the different a_i^d .

$$\lambda(t) = \sum_{i=1}^d a_i^d I_{[b_{i-1}^d, b_i^d]}(t) \quad \text{where } b_i^d = i T_{\max}^d$$

As seen above, we know that the formula for the log likelihood is the following:

$$L_\lambda(N) = \sum_{t \in N} \log(\lambda(t)) - \int_0^{T_{\max}} \lambda(t) dt$$

$$= \int_0^{T_{\max}} \log(\lambda(t)) dN_t - \int_0^{T_{\max}} \lambda(t) dt$$

We know that this is equal to a_i^d as long as t is in between b_{i-1}^d & b_i^d

\Rightarrow We can hence rewrite the formula like this:

$$\sum_{i=1}^d \left((\log a_i^d) N_{[b_{i-1}^d, b_i^d]} - a_i^d (b_i^d - b_{i-1}^d) \right)$$

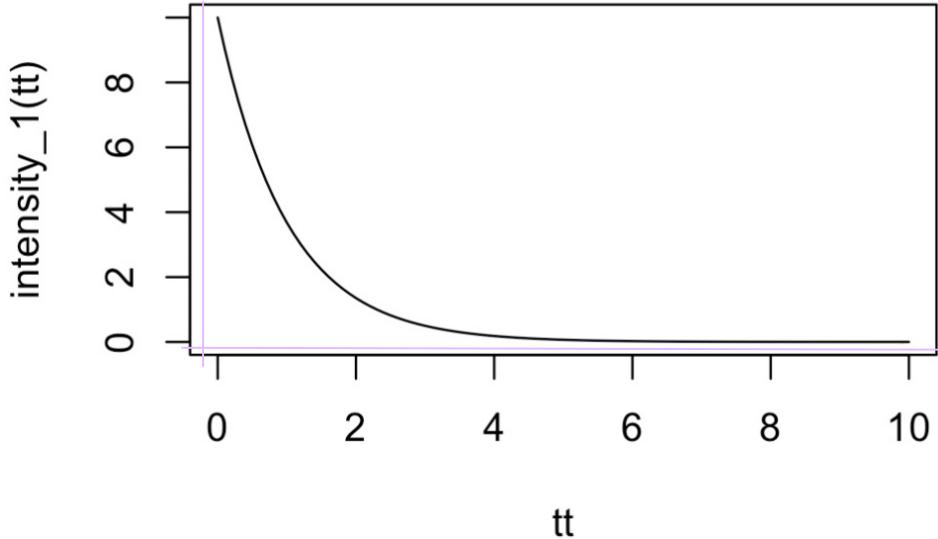
from the ex. we know that $b_i^d = i T_{\max}^d$, so we know that this is T_{\max}^d

\Rightarrow We do derivation again:

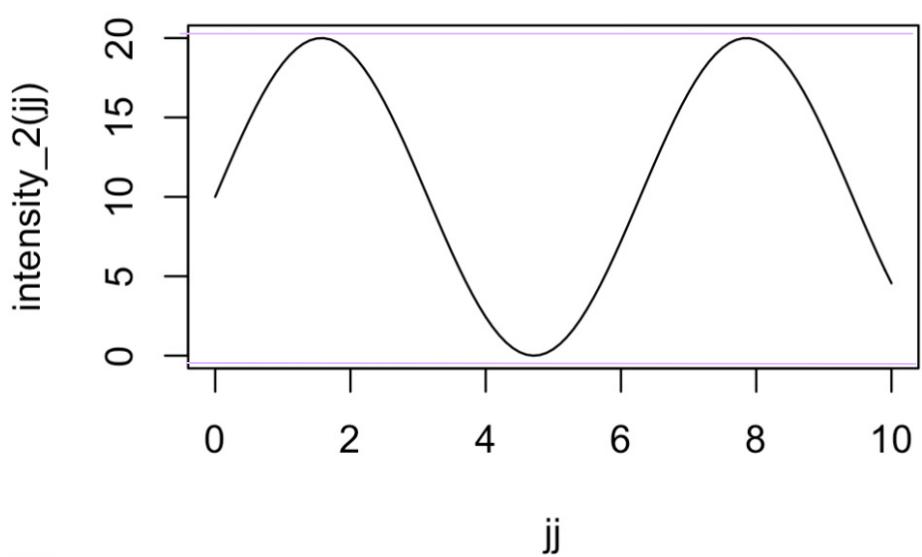
$$\frac{dL}{da_i^d} \quad \& \text{ find that } \hat{a}_i^d = \frac{N_{[b_{i-1}^d, b_i^d]}}{T_{\max}^d}$$

c) Simulate by thinning a Poisson Process of intensity ne^{-x} and $n(1 + \sin(x))$ on $(0, T_{\max})$ for various n .

Superpose different MLE of the intensities for various choices of d & n .

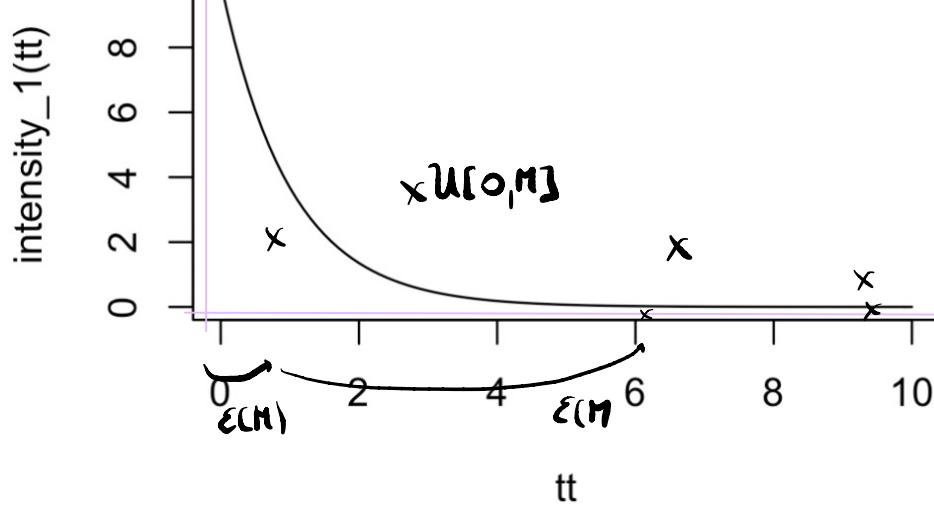


\Rightarrow We can see that the first intensity function is bounded by 0.

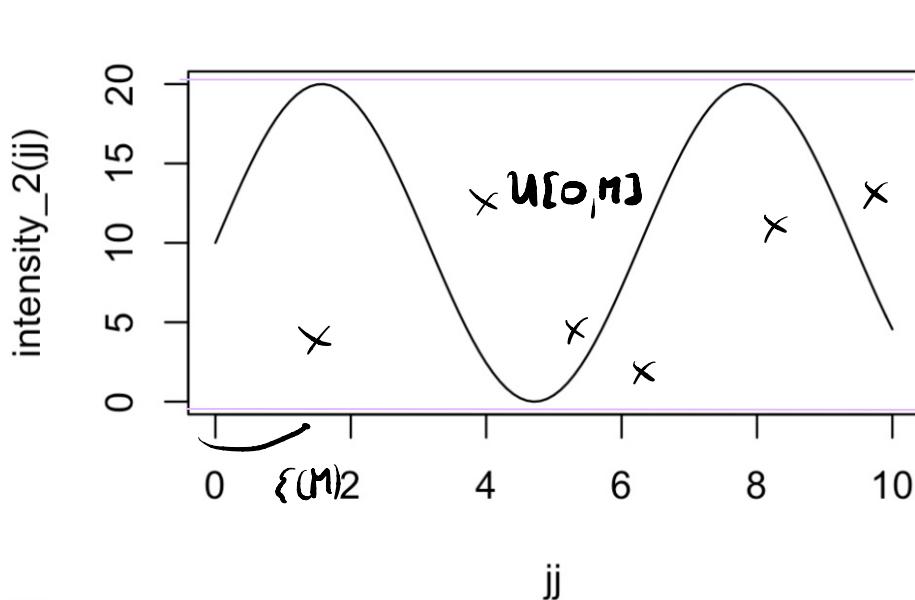


\Rightarrow We can see that the second intensity function is bounded by 20.

Hence, the thinning Process :



where $M=1$



where $M=1$

```
#first intensity function
intensity_1 = function(x, n){
  return(n*exp(-x))
}

#second intensity function
intensity_2 = function(x, n){
  return(n*(1 + sin(x)))
}

intensity_1_plus = 0
intensity_2_plus = 20

Non_Homo_Poisson = function()
{
  t = 0 #t starts out with zero
  T = 1 #t ends with 1
  X = numeric()
  while(t < T){ #we want the process to go on until the cumulative sum of t reaches T, which is 1 here
    u = runif(1) #as long as the above is true, we create u, which is a uniform random variable
    t = t - log(u)/intensity_1_plus #then our new t will be this #it is the same thing as doing it separately and then taking the cumulative sum of it.
    d = runif(1)
    if (d < h(t)/intensity_1_plus){ # compare our uniform variable to the probability
      X = c(X, t) #if the above is the case, we add this to our X
    }
  }
  return(X)
}
```

d) Apply AIC criterion to select for a given observed Poisson Process, a good d.

MISSING

③ How do we check that our simulation is correct?

MISSING